

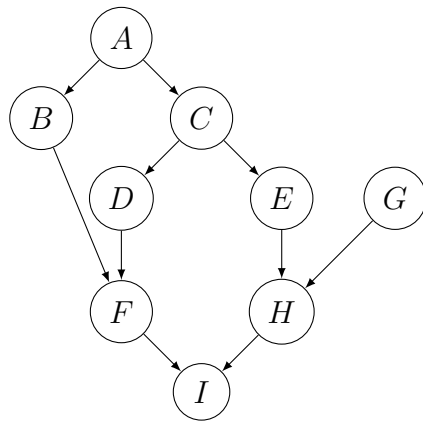
# COMP0085 Summative Assignment

Jan 4, 2023

## Question 1

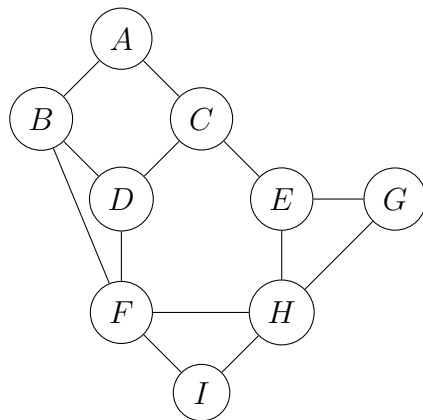
(a)

The directed acyclic graph:

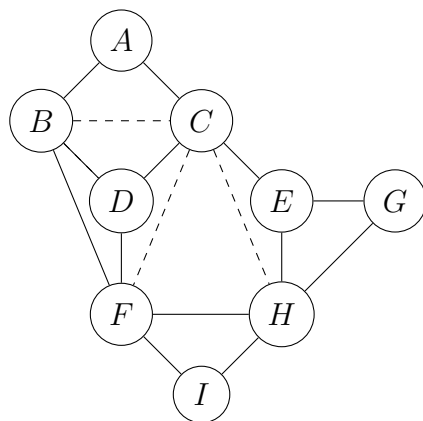


(b)

The moralised graph:

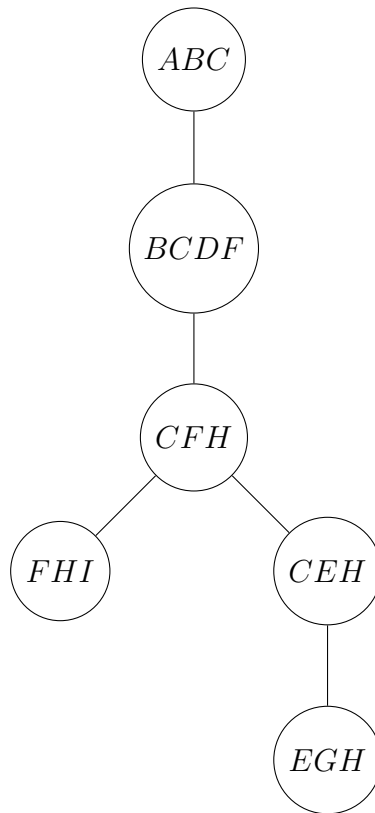


An effective triangulation:



where the dashed lines are edges added to triangulate the moralised graph.

The resulting junction tree:



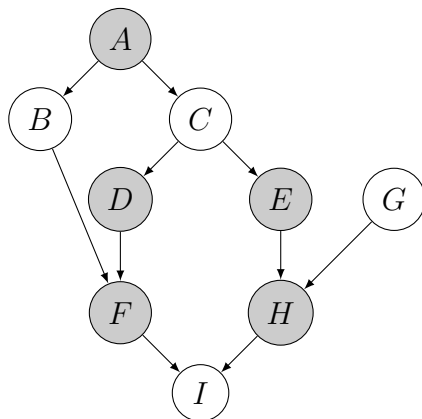
where the circular nodes are cliques.

The junction tree redrawn as a factor graph:



where the circular nodes are cliques and the square nodes are separators/factors.

(c)



The set  $\{A, D, E, F, H\}$  is a non-unique smallest set of molecules such that if the concentrations of the species within the set are known, the concentrations of the others  $\{B, C, G, I\}$  would all be independent (conditioned on the measured ones).

(d)

(e)

## Question 2

(a)

We want the posterior mean and covariance over  $a$  and  $b$ . Defining a weight vector  $\mathbf{w}$ :

$$\mathbf{w} = \begin{bmatrix} a \\ b \end{bmatrix}$$

Our distribution for  $\mathbf{w}$ :

$$P(\mathbf{w}) = \mathcal{N} \left( \begin{bmatrix} \mu_a \\ \mu_b \end{bmatrix}, \begin{bmatrix} \sigma_a^2 & 0 \\ 0 & \sigma_b^2 \end{bmatrix} \right) = \mathcal{N}(\mu_{\mathbf{w}}, \Sigma_{\mathbf{w}})$$

Moreover, for our data  $\mathcal{D} = \{\mathbf{X}, \mathbf{Y}\}$ :

$$P(\mathcal{D}|\mathbf{w}) = \mathcal{N}(\mathbf{Y} - \mathbf{w}^T \mathbf{X}, \sigma^2 \mathbf{I})$$

where  $\mathbf{X} = \begin{bmatrix} t_1 & t_2 \cdots t_N \\ 1 & 1 \cdots 1 \end{bmatrix} \in \mathbb{R}^{2 \times N}$  and  $\mathbf{Y} \in \mathbb{R}^{1 \times N}$ .

Knowing:

$$P(\mathbf{w}|\mathcal{D}) \propto P(\mathcal{D}|\mathbf{w})P(\mathbf{w})$$

we can substitute the above distributions:

$$P(\mathbf{w}|\mathcal{D}) \propto \exp \left( \frac{-1}{2\sigma^2} (\mathbf{Y} - \mathbf{w}^T \mathbf{X}) (\mathbf{Y} - \mathbf{w}^T \mathbf{X})^T \right) \exp \left( \frac{-1}{2} (\mathbf{w} - \mu_{\mathbf{w}})^T \Sigma_{\mathbf{w}}^{-1} (\mathbf{w} - \mu_{\mathbf{w}}) \right)$$

expanding:

$$\log P(\mathbf{w}|\mathcal{D}) \propto \frac{-1}{2} \left( \frac{\mathbf{Y}\mathbf{Y}^T}{\sigma^2} - 2\mathbf{w}^T \frac{\mathbf{X}\mathbf{Y}^T}{\sigma^2} + \mathbf{w}^T \frac{\mathbf{X}\mathbf{X}^T}{\sigma^2} \mathbf{w} + \mathbf{w}^T \Sigma_{\mathbf{w}}^{-1} \mathbf{w} - 2\mathbf{w}^T \Sigma_{\mathbf{w}}^{-1} \mu_{\mathbf{w}} + \mu_{\mathbf{w}}^T \Sigma_{\mathbf{w}}^{-1} \mu_{\mathbf{w}} \right)$$

collecting  $\mathbf{w}$  terms:

$$\log P(\mathbf{w}|\mathcal{D}) \propto \frac{-1}{2} \left( \mathbf{w}^T \left( \frac{\mathbf{X}\mathbf{X}^T}{\sigma^2} + \Sigma_{\mathbf{w}}^{-1} \right) \mathbf{w} - 2\mathbf{w}^T \left( \frac{\mathbf{X}\mathbf{Y}^T}{\sigma^2} + \Sigma_{\mathbf{w}}^{-1} \mu_{\mathbf{w}} \right) \right)$$

Knowing that the posterior  $P(\mathbf{w}|\mathcal{D})$  will be Gaussian with mean  $\bar{\mu}_w$  and covariance  $\bar{\Sigma}_w$ , we can see that expanding the exponent component would have the form:

$$(\mathbf{w} - \bar{\mu}_w)^T \bar{\Sigma}_w^{-1} (\mathbf{w} - \bar{\mu}_w) = \mathbf{w}^T \bar{\Sigma}_w^{-1} \mathbf{w} - 2\mathbf{w}^T \bar{\Sigma}_w^{-1} \bar{\mu}_w + \bar{\mu}_w^T \bar{\Sigma}_w^{-1} \bar{\mu}_w$$

Thus we can identify the posterior covariance:

$$\bar{\Sigma}_w = \left( \frac{\mathbf{X}\mathbf{X}^T}{\sigma^2} + \Sigma_{\mathbf{w}}^{-1} \right)^{-1}$$

and the posterior mean:

$$\bar{\mu}_w = \bar{\Sigma}_w \left( \frac{\mathbf{X}\mathbf{Y}^T}{\sigma^2} + \Sigma_{\mathbf{w}}^{-1} \mu_{\mathbf{w}} \right)$$

## The Python code:

```

1 import scipy
2 import numpy as np
3 from dataclasses import dataclass
4 import pandas as pd
5 import dataframe_image as dfi
6 import matplotlib.pyplot as plt
7
8
9 @dataclass
10 class LinearRegressionParameters:
11     mean: np.ndarray
12     covariance: np.ndarray
13
14     @property
15     def precision(self):
16         return np.linalg.inv(self.covariance)
17
18     def predict(self, x: np.ndarray) -> np.ndarray:
19         return self.mean.T @ x
20
21
22 @dataclass
23 class Theta:
24     linear_regression_parameters: LinearRegressionParameters
25     sigma: float
26
27     @property
28     def variance(self):
29         return self.sigma**2
30
31     @property
32     def precision(self):
33         return 1 / self.variance
34
35
36 def compute_posterior(
37     x: np.ndarray,
38     y: np.ndarray,
39     prior_linear_regression_parameters: LinearRegressionParameters,
40     residuals_precision: float,
41 ) -> LinearRegressionParameters:
42     """
43     Compute the parameters of the posterior distribution on the linear regression weights
44
45     :param x: design matrix (number of features, number of data points)
46     :param y: response matrix (1, number of data points)
47     :param prior_linear_regression_parameters: parameters for the prior distribution on the linear regression
48     :param residuals_precision: the precision of the residuals of the linear regression
49     :return: parameters for the posterior distribution on the linear regression weights
50     """
51     posterior_covariance = np.linalg.inv(
52         residuals_precision * x @ x.T + prior_linear_regression_parameters.precision
53     )
54     posterior_mean = posterior_covariance @ (
55         residuals_precision * x @ y.T
56         + prior_linear_regression_parameters.precision
57         @ prior_linear_regression_parameters.mean
58     )
59     return LinearRegressionParameters(
60         mean=posterior_mean, covariance=posterior_covariance
61     )
62
63
64 def construct_design_matrix(t: np.ndarray):
65     return np.stack((t, np.ones(t.shape)), axis=1).T
66
67
68 def a(
69     t: np.ndarray,
70     y: np.ndarray,
71     sigma: float,
72     prior_linear_regression_parameters: LinearRegressionParameters,
73     save_path: str,
74 ) -> LinearRegressionParameters:
75     x = construct_design_matrix(t)
76     prior_theta = Theta(
77         linear_regression_parameters=prior_linear_regression_parameters,
78         sigma=sigma,
79     )
80     posterior_linear_regression_parameters = compute_posterior(
81         x,
82         y,
83         prior_linear_regression_parameters,
84         residuals_precision=prior_theta.precision,
85     )
86     df_mean = pd.DataFrame(
87         posterior_linear_regression_parameters.mean, columns=["value"]
88     )
89     df_mean.index = ["a", "b"]
90     df_mean = pd.concat([df_mean], keys=["parameters"])
91     dfi.export(df_mean, save_path + "-mean.png")
92
93     df_covariance = pd.DataFrame(

```



```

94     posterior_linear_regression_parameters.covariance, columns=["a", "b"]
95 )
96 df_covariance.index = ["a", "b"]
97 df_covariance = pd.concat([df_covariance], keys=["parameters"])
98 df_covariance = pd.concat([df_covariance.T], keys=["parameters"])
99 dfi.export(df_covariance, save_path + "-covariance.png")
100 return posterior_linear_regression_parameters
101
102
103 def b(
104     t_year, t, y, linear_regression_parameters, error_mean, error_variance, save_path
105 ):
106     x = construct_design_matrix(t)
107     residuals = y - linear_regression_parameters.predict(x)
108     plt.plot(t_year.reshape(-1), residuals.reshape(-1))
109     plt.xlabel("date (decimal year)")
110     plt.ylabel("residual")
111     plt.title("2b: g_obs(t)")
112     plt.savefig(save_path + "-residuals-timeseries")
113     plt.close()
114
115     count, bins = np.histogram(residuals, bins=100, density=True)
116     plt.bar(bins[1:], count, label="residuals")
117     plt.plot(
118         bins[1:],
119         scipy.stats.norm.pdf(bins[1:], loc=error_mean, scale=error_variance),
120         color="red",
121         label="e(t)",
122     )
123     plt.xlabel("residual bin")
124     plt.ylabel("density")
125     plt.title("2b: Residuals Density")
126     plt.legend()
127     plt.savefig(save_path + "-residuals-density-estimation")
128     plt.close()

```

src/solutions/q2.py

### Question 3

(a)

The free energy is can be calculated as:

$$\mathcal{F}(q, \theta) = \langle \log P(\mathbf{x}, \mathbf{s}|\theta) \rangle_{q(\mathbf{s})} + H[Q(\mathbf{s})]$$

Knowing,

$$\log P(\mathbf{x}, \mathbf{s}|\theta) = \log P(\mathbf{x}|\mathbf{s}, \theta) + \log P(\mathbf{s}|\theta)$$

we can write:

$$\mathcal{F}(Q, \theta) = \langle \log P(\mathbf{x}|\mathbf{s}, \theta) \rangle_{q(\mathbf{s})} + \langle \log P(\mathbf{s}|\theta) \rangle_{q(\mathbf{s})} + H[q(\mathbf{s})]$$

Moreover, our mean field approximation:

$$q(\mathbf{s}) = \prod_{i=1}^K q_i(s_i)$$

where  $q_i(s_i) = \lambda_i^{s_i} (1 - \lambda_i)^{(1-s_i)}$ .

To compute the first term:

$$P(\mathbf{x}|\mathbf{s}, \theta) = \mathcal{N} \left( \sum_{i=1}^K s_i \mu_i, \sigma^2 \mathbf{I} \right)$$

substituting the appropriate terms:

$$P(\mathbf{x}|\mathbf{s}, \theta) = 2\pi^{-\frac{d}{2}} |\sigma^2 \mathbf{I}|^{-\frac{1}{2}} \exp \left( -\frac{1}{2} \left( \mathbf{x} - \sum_{i=1}^K s_i \mu_i \right)^T \frac{1}{\sigma^2} \mathbf{I} \left( \mathbf{x} - \sum_{i=1}^K s_i \mu_i \right) \right)$$

with  $d$  being the number of dimensions.

Taking the logarithm:

$$\log P(\mathbf{x}|\mathbf{s}, \theta) = -\frac{d}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \left( \mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K s_i \mu_i + \sum_{i=1}^K \sum_{j=1}^K s_i s_j \mu_i^T \mu_j \right)$$

The expectation distributed to the relevant terms:

$$\langle \log P(\mathbf{x}|\mathbf{s}, \theta) \rangle_{q(\mathbf{s})} = -\frac{d}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \left( \mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K \langle s_i \rangle_{q_i(s_i)} \mu_i + \sum_{i=1}^K \sum_{j=1}^K \langle s_i s_j \rangle_{q_i(s_i) q_j(s_j)} \mu_i^T \mu_j \right)$$

Evaluating the expectations:

$$\langle \log P(\mathbf{x}|\mathbf{s}, \theta) \rangle_{q(\mathbf{s})} = -\frac{d}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \left( \mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K \lambda_i \mu_i + \sum_{i=1}^K \sum_{j=1, j \neq i}^K \lambda_i \lambda_j \mu_i^T \mu_j + \sum_{i=1}^K \lambda_i \mu_i^T \mu_i \right)$$

where  $\langle s_i s_j \rangle_{q_i(s_i)} = \langle s_i \rangle_{q_i(s_i)}$  because  $s_i \in \{0, 1\}$ .

To compute the second term:

$$P(\mathbf{s}|\theta) = \prod_{i=1}^K \pi_i^{s_i} (1 - \pi_i)^{(1-s_i)}$$

Taking the logarithm:

$$\log P(\mathbf{s}|\theta) = \sum_{i=1}^K s_i \log \pi_i + (1 - s_i) \log(1 - \pi_i)$$

The expectation distributed to the relevant terms:

$$\langle \log P(\mathbf{s}|\theta) \rangle_{q(\mathbf{s})} = \sum_{i=1}^K \langle s_i \rangle_{q_i(s_i)} \log \pi_i + (1 - \langle s_i \rangle_{q_i(s_i)}) \log(1 - \pi_i)$$

Evaluating the expectations:

$$\langle \log P(\mathbf{s}|\theta) \rangle_{q(\mathbf{s})} = \sum_{i=1}^K \lambda_i \log \pi_i + (1 - \lambda_i) \log(1 - \pi_i)$$

To compute the third term, we use the mean field factorisation:

$$H[q(\mathbf{s})] = \sum_{i=1}^K H[q_i(s_i)]$$

Thus,

$$H[q(\mathbf{s})] = - \sum_{i=1}^K \sum_{s_i \in \{0,1\}} q_i(s_i) \log q_i(s_i)$$

Substituting the appropriate values:

$$H[q(\mathbf{s})] = - \sum_{i=1}^K \lambda_i \log \lambda_i + (1 - \lambda_i) \log(1 - \lambda_i)$$

Combining, we have our free energy expression:

$$\begin{aligned} \mathcal{F}(q, \theta) = & \frac{-d}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \left( \mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K \lambda_i \mu_i + \sum_{i=1}^K \sum_{j=1, j \neq i}^K \lambda_i \lambda_j \mu_i^T \mu_j + \sum_{i=1}^K \lambda_i \mu_i^T \mu_i \right) \\ & + \sum_{i=1}^K \lambda_i \log \pi_i + (1 - \lambda_i) \log(1 - \pi_i) \\ & - \sum_{i=1}^K \lambda_i \log \lambda_i + (1 - \lambda_i) \log(1 - \lambda_i) \end{aligned}$$

To derive the partial update for  $q_i(s_i)$  we take the variational derivative of the Lagrangian, enforcing the normalisation of  $q_i$ :

$$\frac{\partial}{\partial q_i} \left( \mathcal{F}(q, \theta) + \lambda^{LG} \int q_i - 1 \right) = \langle \log P(\mathbf{x}, \mathbf{s}|\theta) \rangle_{\prod_{j \neq i} q_j(s_j)} - \log q_i(s_i) - 1 + \lambda^{LG}$$

where  $\lambda^{LG}$  is the Lagrange multiplier.

Setting this to zero we can solve for the  $\lambda_i$  that maximises the free energy:

$$\log q_i(s_i) = \langle \log P(\mathbf{x}, \mathbf{s} | \theta) \rangle_{\prod_{j \neq i} q_j(s_j)} - 1 + \lambda^{LG}$$

Similar to our free energy derivation:

$$\langle \log P(\mathbf{x} | \mathbf{s}, \theta) \rangle_{\prod_{j \neq i} q_j(s_j)} \propto -\frac{1}{2\sigma^2} \left( \mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{k=1}^K \langle s_k \rangle_{\prod_{j \neq i} q_j(s_j)} \mu_k + \sum_{k=1}^K \sum_{j=1}^K \langle s_k s_j \rangle_{\prod_{j \neq i} q_j(s_j)} \right)$$

and

$$\langle \log P(\mathbf{s} | \theta) \rangle_{\prod_{j \neq i} q_j(s_j)} = \sum_{k=1}^K \langle s_k \rangle_{\prod_{j \neq i} q_j(s_j)} \log \pi_k + (1 - \langle s_k \rangle_{\prod_{j \neq i} q_j(s_j)}) \log(1 - \pi_k)$$

We can write:

$$\log q_i(s_i) \propto \log P(\mathbf{x} | \mathbf{s}, \theta)_{\prod_{j \neq i} q_j(s_j)} + \langle \log P(\mathbf{s} | \theta) \rangle_{\prod_{j \neq i} q_j(s_j)}$$

Substituting the relevant terms:

$$\log q_i(s_i) \propto -\frac{1}{2\sigma^2} \left( -2s_i \mathbf{x}^T \mu_i + s_i s_i \mu_i^T \mu_i + 2 \sum_{j=1, j \neq i}^K s_i \lambda_j \mu_i^T \mu_j \right) + s_i \log \pi_i + (1 - s_i) \log(1 - \pi_i)$$

Knowing  $\log q_i(s_i) = s_i \log \lambda_i + (1 - s_i) \log(1 - \lambda_i)$ :

$$\log q_i(s_i) \propto s_i \log \frac{\lambda_i}{1 - \lambda_i}$$

Thus,

$$s_i \log \frac{\lambda_i}{1 - \lambda_i} \propto -\frac{1}{2\sigma^2} \left( -2s_i \mathbf{x}^T \mu_i + s_i s_i \mu_i^T \mu_i + 2 \sum_{j=1, j \neq i}^K s_i \lambda_j \mu_i^T \mu_j \right) + s_i \log \frac{\pi_i}{1 - \pi_i}$$

Also, because  $s_i \in \{0, 1\}$  we know that  $s_i^2 = s_i$ :

$$s_i \log \frac{\lambda_i}{1 - \lambda_i} \propto -\frac{1}{2\sigma^2} \left( -2s_i \mathbf{x}^T \mu_i + s_i \mu_i^T \mu_i + 2 \sum_{j=1, j \neq i}^K s_i \lambda_j \mu_i^T \mu_j \right) + s_i \log \frac{\pi_i}{1 - \pi_i}$$

Because we have only kept terms with  $s_i$ , this is an equality:

$$s_i \log \frac{\lambda_i}{1 - \lambda_i} = \frac{s_i \mu_i^T}{2\sigma^2} \left( 2\mathbf{x} - \mu_i - 2 \sum_{j=1, j \neq i}^K \lambda_j \mu_j \right) + s_i \log \frac{\pi_i}{1 - \pi_i}$$

Solving for  $\lambda_i$ :

$$\lambda_i = \frac{1}{1 + \exp \left[ - \left( \frac{\mu_i^T}{\sigma^2} \left( \mathbf{x} - \frac{\mu_i}{2} - \sum_{j=1, j \neq i}^K \lambda_j \mu_j \right) + \log \frac{\pi_i}{1 - \pi_i} \right) \right]}$$

we have our partial update.

## The Python code:

```
1 from dataclasses import dataclass
2
3 import numpy as np
4
5 from demo_code.MStep import m_step
6
7
8 @dataclass
9 class MeanFieldApproximation:
10     """
11     lambda_matrix: parameters variational approximation (number_of_points, number_of_latent_variables)
12     """
13
14     lambda_matrix: np.ndarray
15
16     def lambda_matrix_exclude(self, exclude_latent_index: int) -> np.ndarray:
17         # (number_of_points, number_of_latent_variables-1)
18         return np.concatenate(
19             (
20                 self.lambda_matrix[:, :exclude_latent_index],
21                 self.lambda_matrix[:, exclude_latent_index + 1 :],
22             ),
23             axis=1,
24         )
25
26     @property
27     def log_lambda_matrix(self):
28         return np.log(self.lambda_matrix)
29
30     @property
31     def log_one_minus_lambda_matrix(self):
32         return np.log(1 - self.lambda_matrix)
33
34     @property
35     def n(self):
36         return self.lambda_matrix.shape[0]
37
38     @property
39     def k(self):
40         return self.lambda_matrix.shape[1]
41
42
43 def init_mean_field_approximation(k: int, n: int) -> MeanFieldApproximation:
44     return MeanFieldApproximation(
45         lambda_matrix=np.random.random(size=(n, k)),
46     )
47
48
49 @dataclass
50 class BinaryLatentFactorModel:
51     """
52     mu: matrix of means (number_of_dimensions, number_of_latent_variables)
53     sigma: gaussian noise parameter
54     pi: vector of priors (1, number_of_latent_variables)
55     """
56
57     mu: np.ndarray
58     sigma: float
59     pi: np.ndarray
60
61     def mu_exclude(self, exclude_latent_index: int) -> np.ndarray:
62         # (number_of_dimensions, number_of_latent_variables-1)
63         return np.concatenate(
64             (self.mu[:, :exclude_latent_index], self.mu[:, exclude_latent_index + 1 :]),
65             axis=1,
66         )
67
68     @property
69     def log_pi(self):
70         return np.log(self.pi)
71
72     @property
73     def log_one_minus_pi(self):
74         return np.log(1 - self.pi)
75
76     @property
77     def variance(self):
78         return self.sigma**2
79
80     @property
81     def precision(self):
82         return 1 / self.variance
83
84     @property
85     def d(self):
86         return self.mu.shape[0]
87
88     @property
89     def k(self):
90         return self.mu.shape[1]
91
92
93 def init_binary_latent_factor_model(
94     x: np.ndarray,
```

```

95     mean_field_approximation: MeanFieldApproximation,
96 ) -> BinaryLatentFactorModel:
97     return maximisation_step(x, mean_field_approximation)
98
99
100 def _compute_expectation_log_p_x_s_given_theta(
101     x: np.ndarray,
102     binary_latent_factor_model: BinaryLatentFactorModel,
103     mean_field_approximation: MeanFieldApproximation,
104 ) -> float:
105     """
106     The first term of the free energy, the expectation of log P(X,S|theta)
107
108     :param x: data matrix (number_of_points, number_of_dimensions)
109     :param binary_latent_factor_model: a binary_latent_factor_model
110     :param mean_field_approximation: a mean_field_approximation
111     :return: the expectation of log P(X,S|theta)
112     """
113     # (number_of_points, number_of_dimensions)
114     mu_lambda = mean_field_approximation.lambda_matrix @ binary_latent_factor_model.mu.T
115
116     # (number_of_latent_variables, number_of_latent_variables)
117     expectation_s_i_s_j_mu_i_mu_j = np.multiply(
118         mean_field_approximation.lambda_matrix.T
119         @ mean_field_approximation.lambda_matrix,
120         binary_latent_factor_model.mu.T @ binary_latent_factor_model.mu,
121     )
122
123     term1 = mean_field_approximation.n*(
124         (-binary_latent_factor_model.d / 2) * np.log(2 * np.pi*binary_latent_factor_model.variance)
125     )
126
127     term2 = - (0.5 * binary_latent_factor_model.precision) * (
128         np.sum(np.multiply(x, x))
129         - 2 * np.sum(np.multiply(x, mu_lambda))
130         + np.sum(expectation_s_i_s_j_mu_i_mu_j)
131         - np.trace(
132             expectation_s_i_s_j_mu_i_mu_j
133         ) # remove incorrect E[s_i s_i] = lambda_i * lambda_i
134         + np.sum( # add correct E[s_i s_i] = lambda_i
135             mean_field_approximation.lambda_matrix
136             @ np.multiply(
137                 binary_latent_factor_model.mu, binary_latent_factor_model.mu
138             ).T
139         )
140     )
141
142     expectation_log_p_x_given_s_theta = (
143         mean_field_approximation.n * (
144             (-binary_latent_factor_model.d / 2) * np.log(2 * np.pi * binary_latent_factor_model.variance)
145         )
146         - (0.5 * binary_latent_factor_model.precision) * (
147             np.sum(np.multiply(x, x))
148             - 2 * np.sum(np.multiply(x, mu_lambda))
149             + np.sum(expectation_s_i_s_j_mu_i_mu_j)
150             - np.trace(
151                 expectation_s_i_s_j_mu_i_mu_j
152             ) # remove incorrect E[s_i s_i] = lambda_i * lambda_i
153             + np.sum( # add correct E[s_i s_i] = lambda_i
154                 mean_field_approximation.lambda_matrix
155                 @ np.multiply(
156                     binary_latent_factor_model.mu, binary_latent_factor_model.mu
157                 ).T
158             )
159         )
160     )
161
162     expectation_log_p_s_given_theta = np.sum(
163         np.multiply(
164             mean_field_approximation.lambda_matrix,
165             binary_latent_factor_model.log_pi,
166         )
167         + np.multiply(
168             1 - mean_field_approximation.lambda_matrix,
169             binary_latent_factor_model.log_one_minus_pi,
170         )
171     )
172     return expectation_log_p_x_given_s_theta + expectation_log_p_s_given_theta
173
174 def _compute_mean_field_approximation_entropy(
175     mean_field_approximation: MeanFieldApproximation,
176 ) -> float:
177     return -np.sum(
178         np.multiply(
179             mean_field_approximation.lambda_matrix,
180             mean_field_approximation.log_lambda_matrix,
181         )
182         + np.multiply(
183             1 - mean_field_approximation.lambda_matrix,
184             mean_field_approximation.log_one_minus_lambda_matrix,
185         )
186     )
187
188
189 def compute_free_energy(
190     x: np.ndarray,

```

```

191     binary_latent_factor_model: BinaryLatentFactorModel,
192     mean_field_approximation: MeanFieldApproximation,
193 ) -> float:
194     """
195     free energy associated with current EM parameters and data x
196
197     :param x: data matrix (number_of_points, number_of_dimensions)
198     :param binary_latent_factor_model: a binary_latent_factor_model
199     :param mean_field_approximation: a mean_field_approximation
200     :return: free energy
201     """
202     expectation_log_p_x_s_given_theta = _compute_expectation_log_p_x_s_given_theta(
203         x, binary_latent_factor_model, mean_field_approximation
204     )
205     mean_field_approximation_entropy = _compute_mean_field_approximation_entropy(
206         mean_field_approximation
207     )
208     return expectation_log_p_x_s_given_theta + mean_field_approximation_entropy
209
210
211 def partial_expectation_step(
212     x: np.ndarray,
213     binary_latent_factor_model: BinaryLatentFactorModel,
214     mean_field_approximation: MeanFieldApproximation,
215     latent_factor: int,
216 ) -> np.ndarray:
217     """ Partial Variational E step for factor i for all data points
218
219     :param x: data matrix (number_of_points, number_of_dimensions)
220     :param binary_latent_factor_model: a binary_latent_factor_model
221     :param mean_field_approximation: a mean_field_approximation
222     :param latent_factor: latent factor to compute partial update
223     :return: lambda_vector: new lambda parameters for the latent factor (number_of_points, 1)
224     """
225     lambda_matrix_excluded = mean_field_approximation.lambda_matrix_exclude(
226         latent_factor
227     )
228     mu_excluded = binary_latent_factor_model.mu_exclude(latent_factor)
229
230     mu_latent = binary_latent_factor_model.mu[:, latent_factor]
231     # (number_of_points, 1)
232     partial_expectation_log_p_x_given_s_theta_proportion = (
233         binary_latent_factor_model.precision
234         * (
235             x # (number_of_points, number_of_dimensions)
236             - 0.5 * mu_latent.T # (1, number_of_dimensions)
237             - lambda_matrix_excluded # (number_of_points, number_of_latent_variables-1)
238             @ mu_excluded.T # (number_of_latent_variables-1, number_of_dimensions)
239         )
240         @ mu_latent # (number_of_dimensions, 1)
241     )
242
243     # (1, 1)
244     partial_expectation_log_p_s_given_theta_proportion = np.log(
245         binary_latent_factor_model.pi[0, latent_factor]
246         / (1 - binary_latent_factor_model.pi[0, latent_factor])
247     )
248
249     # (number_of_points, 1)
250     partial_expectation_log_p_x_s_given_theta_proportion = (
251         partial_expectation_log_p_x_given_s_theta_proportion
252         + partial_expectation_log_p_s_given_theta_proportion
253     )
254
255     # (number_of_points, 1)
256     lambda_vector = 1 / (
257         1 + np.exp(-partial_expectation_log_p_x_s_given_theta_proportion)
258     )
259     lambda_vector[lambda_vector == 0] = 1e-10
260     lambda_vector[lambda_vector == 1] = 1 - 1e-10
261     return lambda_vector
262
263
264 def variational_expectation_step(
265     x: np.ndarray,
266     binary_latent_factor_model: BinaryLatentFactorModel,
267     mean_field_approximation: MeanFieldApproximation,
268     max_steps: int,
269     convergence_criterion: float,
270 ) -> MeanFieldApproximation:
271     """ Variational E step
272
273     :param x: data matrix (number_of_points, number_of_dimensions)
274     :param binary_latent_factor_model: a binary_latent_factor_model
275     :param mean_field_approximation: a mean_field_approximation
276     :param max_steps: maximum number of steps of fixed point equations
277     :param convergence_criterion: early stopping if change in free energy < convergence_criterion
278     :return: mean field approximation
279     """
280     # previous_free_energy = compute_free_energy(
281     #     x, binary_latent_factor_model, mean_field_approximation
282     # )
283     for i in range(max_steps):
284         for latent_factor in range(binary_latent_factor_model.k):
285             mean_field_approximation.lambda_matrix[
286                 :, latent_factor

```

```

287         ] = partial_expectation_step(
288             x, binary_latent_factor_model, mean_field_approximation, latent_factor
289         )
290         # free_energy = compute_free_energy(
291         #     x, binary_latent_factor_model, mean_field_approximation
292         # )
293         # if free_energy - previous_free_energy <= convergence_criterion:
294         #     break
295         # previous_free_energy = free_energy
296     return mean_field_approximation
297
298
299 def maximisation_step(
300     x: np.ndarray,
301     mean_field_approximation: MeanFieldApproximation,
302 ) -> BinaryLatentFactorModel:
303     expectation_s = mean_field_approximation.lambda_matrix
304     expectation_ss = (
305         mean_field_approximation.lambda_matrix.T
306         @ mean_field_approximation.lambda_matrix
307     )
308     np.fill_diagonal(expectation_ss, mean_field_approximation.lambda_matrix.sum(axis=0))
309     mu, sigma, pi = m_step(x, expectation_s, expectation_ss)
310     return BinaryLatentFactorModel(
311         mu=mu,
312         sigma=sigma,
313         pi=pi,
314     )
315
316
317 def learn_binary_factors(
318     x: np.ndarray,
319     k: int,
320     em_maximum_iterations: int,
321     e_maximum_steps: int,
322     e_convergence_criterion: float,
323 ):
324     n = x.shape[0]
325     mean_field_approximation = init_mean_field_approximation(k, n)
326     binary_latent_factor_model = init_binary_latent_factor_model(
327         x, mean_field_approximation
328     )
329
330     for _ in range(em_maximum_iterations):
331         mean_field_approximation = variational_expectation_step(
332             x=x,
333             binary_latent_factor_model=binary_latent_factor_model,
334             mean_field_approximation=mean_field_approximation,
335             max_steps=e_maximum_steps,
336             convergence_criterion=e_convergence_criterion,
337         )
338         free_energy = compute_free_energy(
339             x, binary_latent_factor_model, mean_field_approximation
340         )
341         binary_latent_factor_model = maximisation_step(
342             x=x,
343             mean_field_approximation=mean_field_approximation,
344         )
345         free_energy = compute_free_energy(
346             x, binary_latent_factor_model, mean_field_approximation
347         )
348     return mean_field_approximation, binary_latent_factor_model

```

src/solutions/q3.py



## Question 4

## Question 5

(a)

The log-joint probability for a single observation-source pair:

$$\log p(\mathbf{s}, \mathbf{x}) = \log p(\mathbf{s}) + (\mathbf{x}|\mathbf{s})$$

Knowing  $p(\mathbf{s}) = \prod_{i=1}^K p(s_i|\pi_i)$  and  $p(\mathbf{x}|\mathbf{s}) = \mathcal{N}(\sum_{i=1}^K s_i \mu_i, \sigma^2 \mathbf{I})$ :

$$\log p(\mathbf{s}, \mathbf{x}) \propto \frac{-1}{2} \left( \mathbf{x} - \sum_{i=1}^K s_i \mu_i \right)^T \frac{1}{\sigma^2} \mathbf{I} \left( \mathbf{x} - \sum_{i=1}^K s_i \mu_i \right) + \sum_{i=1}^K (s_i \log \pi_i + (1 - s_i) \log(1 - \pi_i))$$

Expanding:

$$\log p(\mathbf{s}, \mathbf{x}) \propto \frac{-1}{2\sigma^2} \left( \mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K s_i \mu_i + \sum_{i=1}^K \sum_{j=1}^K s_i s_j \mu_i^T \mu_j \right) + \sum_{i=1}^K (s_i \log \pi_i + (1 - s_i) \log(1 - \pi_i))$$

Collecting terms pertaining to  $s_i$ :

$$\log p(\mathbf{s}, \mathbf{x}) = \sum_{i=1}^K \left( \left( \frac{\mathbf{x}^T \mu_i}{\sigma^2} + \log \frac{\pi_i}{1 - \pi_i} \right) s_i \right) + \sum_{i=1}^K \sum_{j=1}^K \left( \frac{-\mu_i^T \mu_j}{2\sigma^2} s_i s_j \right) + C$$

where  $C$  are all other terms without  $s_i$ .

Knowing that  $s_i^2 = s_i$ :

$$\log p(\mathbf{s}, \mathbf{x}) = \sum_{i=1}^K \left( \left( \frac{\mathbf{x}^T \mu_i}{\sigma^2} + \log \frac{\pi_i}{1 - \pi_i} - \frac{\mu_i^T \mu_j}{2\sigma^2} \right) s_i \right) + \sum_{i=1}^K \sum_{j=1}^{i-1} \left( \frac{-\mu_i^T \mu_j}{\sigma^2} s_i s_j \right) + C$$

Thus:

$$\log p(\mathbf{s}, \mathbf{x}) = \sum_{i=1}^K \log f_i(s_i) + \sum_{i=1}^K \sum_{j=1}^{i-1} \log g_{ij}(s_i, s_j)$$

where the factors are defined:

$$\log f_i(s_i) = \left( \frac{\mathbf{x}^T \mu_i}{\sigma^2} + \log \frac{\pi_i}{1 - \pi_i} - \frac{\mu_i^T \mu_j}{2\sigma^2} \right) s_i$$

and

$$\log g_{ij}(s_i, s_j) = \frac{-\mu_i^T \mu_j}{\sigma^2} s_i s_j$$

as required. Note that  $C$  can simply be absorbed into any one of these factors.

The Boltzmann Machine can be defined:

$$P(\mathbf{s}|\mathbf{W}, \mathbf{b}) = \frac{1}{Z} \exp \left( \sum_{i=1}^K \sum_{j=1}^{i-1} W_{ij} s_i s_j - \sum_{i=1}^K b_i s_i \right)$$

where  $s_i \in \{0, 1\}$ , the same as our source variables.

From our factorisation, we can see that  $p(\mathbf{s}, \mathbf{x})$  is a Boltzmann Machine with:

$$W_{ij} = \frac{-\mu_i^T \mu_j}{\sigma^2}$$

and

$$b_i = - \left( \frac{\mathbf{x}^T \mu_i}{\sigma^2} + \log \frac{\pi_i}{1 - \pi_i} - \frac{\mu_i^T \mu_j}{2\sigma^2} \right)$$

and

$$\log Z = -C$$

**(b)**

For  $f_i(s_i)$ , we will choose a Bernoulli approximation:

$$\tilde{f}_i(s_i) = \lambda_i^{s_i} + (1 - \lambda_i)^{1-s_i}$$

Thus,

$$\log \tilde{f}_i(s_i) \propto \log \left( \frac{\lambda_i}{1 - \lambda_i} \right) s_i$$

For  $g_{ij}(s_i, s_j)$ , we will choose a product of Bernoulli's approximation:

$$\tilde{g}_{ij}(s_i, s_j) = (\theta_i^{s_i} + (1 - \theta_i)^{1-s_i}) (\theta_j^{s_j} + (1 - \theta_j)^{1-s_j})$$

Thus,

$$\log \tilde{g}_{ij}(s_i, s_j) \propto \log \left( \frac{\theta_i}{1 - \theta_i} \right) s_i + \log \left( \frac{\theta_j}{1 - \theta_j} \right) s_j$$

To derive the a message passing scheme, we first define:

$$q(\mathbf{s}) = \left( \prod_{i=1}^K \tilde{f}_i(s_i) \right) \left( \prod_{i=1}^K \prod_{j=1}^{i-1} \tilde{g}_{ij}(s_i, s_j) \right)$$

Thus, we can derive cavity distributions:

$$q_{-\tilde{f}_i(s_i)}(\mathbf{s}) = \left( \prod_{i'=1, i' \neq i}^K \tilde{f}_{i'}(s_{i'}) \right) \left( \prod_{i=1}^K \prod_{j=1}^{i-1} \tilde{g}_{ij}(s_i, s_j) \right)$$

and

$$q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) = \left( \prod_{i'=1}^K \tilde{f}_{i'}(s_{i'}) \right) \left( \prod_{i'=1}^K \prod_{\substack{j'=1 \\ i', j' \neq i, j}}^{i'-1} \tilde{g}_{i'j'}(s_{i'}, s_{j'}) \right)$$

For  $\tilde{f}_i(s_i)$ , we do not need to make an approximation step. This is because we are minimising:

$$\tilde{f}_i(s_i) = \arg \min \mathbf{KL} \left[ f_i(s_i) q_{-\tilde{f}_i(s_i)}(\mathbf{s}) \| \tilde{f}_i(s_i) q_{-\tilde{f}_i(s_i)}(\mathbf{s}) \right]$$

We know that the factor  $\log f_i(s_i)$  is a Bernoulli of the form  $b_i s_i$ . Because our approximation is also Bernoulli, we can simply solve for  $\lambda_i$  in  $\log \tilde{f}_i(s_i)$ :

$$\log \tilde{f}_i(s_i) = \log f_i(s_i)$$

$$\log \left( \frac{\lambda_i}{1 - \lambda_i} \right) s_i = b_i s_i$$

$$\lambda_i = \frac{1}{1 + \exp(-b_i)}$$

On the other hand, for  $\tilde{g}_{ij}(s_i, s_j)$ , we will approximate with:

$$\tilde{g}_{ij}(s_i, s_j) = \arg \min \mathbf{KL} \left[ g_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) \| \tilde{g}_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) \right]$$

Note that because  $\tilde{g}_{ij}(s_i, s_j)$  is the product of two Bernoulli distributions, we only require the natural parameters:

$$\phi_{ij}(\theta) = \begin{bmatrix} \theta_i \\ \theta_j \end{bmatrix}$$

the mean with respect to  $s_i$  and  $s_j$  respectively.

We can write:

$$\begin{aligned} \log \tilde{g}_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) &\propto \log \left( \frac{\theta_i}{1 - \theta_i} \right) s_i + \log \left( \frac{\theta_j}{1 - \theta_j} \right) s_j \\ &\quad + \sum_{i'=1}^K \log \left( \frac{\lambda_{i'}}{1 - \lambda_{i'}} \right) s_{i'} \\ &\quad + \sum_{i'=1}^K \sum_{\substack{j'=1 \\ j' \neq i, j}}^{i'-1} \log \left( \frac{\theta_{i'}}{1 - \theta_{i'}} \right) s_{i'} + \log \left( \frac{\theta_{j'}}{1 - \theta_{j'}} \right) s_{j'} \end{aligned}$$

Only including terms with  $s_i$  and  $s_j$ :

$$\begin{aligned} \log \tilde{g}_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) &\propto \left( (K - 1) \log \left( \frac{\theta_i}{1 - \theta_i} \right) + \log \left( \frac{\lambda_i}{1 - \lambda_i} \right) \right) s_i \\ &\quad + \left( (K - 1) \log \left( \frac{\theta_j}{1 - \theta_j} \right) + \log \left( \frac{\lambda_j}{1 - \lambda_j} \right) \right) s_j \end{aligned}$$

Moreover:

$$\begin{aligned} \log g_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) &\propto \frac{-\mu_i^T \mu_j}{\sigma^2} s_i s_j \\ &\quad + \sum_{i'=1}^K \log \left( \frac{\lambda_{i'}}{1 - \lambda_{i'}} \right) s_{i'} \\ &\quad + \sum_{i'=1}^K \sum_{\substack{j'=1 \\ j' \neq i, j}}^{i'-1} \log \left( \frac{\theta_{i'}}{1 - \theta_{i'}} \right) s_{i'} + \log \left( \frac{\theta_{j'}}{1 - \theta_{j'}} \right) s_{j'} \end{aligned}$$

Only including terms with  $s_i$  and  $s_j$ :

$$\begin{aligned} \log g_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) &\propto \frac{-\mu_i^T \mu_j}{\sigma^2} s_i s_j \\ &\quad + \left( (K - 2) \log \left( \frac{\theta_i}{1 - \theta_i} \right) + \log \left( \frac{\lambda_i}{1 - \lambda_i} \right) \right) s_i \\ &\quad + \left( (K - 2) \log \left( \frac{\theta_j}{1 - \theta_j} \right) + \log \left( \frac{\lambda_j}{1 - \lambda_j} \right) \right) s_j \end{aligned}$$

## Appendix 1: constants.py

```
1 import os
2
3 DATA_FOLDER = "data"
4
5 CO2_FILE_PATH = os.path.join(DATA_FOLDER, "co2.txt")
6 IMAGES_FILE_PATH = os.path.join(DATA_FOLDER, "images.jpg")
7
8 OUTPUTS_FOLDER = "outputs"
9
10 DEFAULT_SEED = 0
11
12 M1 = [0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0]
13
14 M2 = [0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0]
15
16 M3 = [1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
17
18 M4 = [1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1]
19
20 M5 = [0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0]
21
22 M6 = [1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1]
23
24 M7 = [0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0]
25
26 M8 = [0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1]
```

src/constants.py

## Appendix 2: main.py

```
1 import os
2 import pandas as pd
3 import numpy as np
4 from src.constants import CO2_FILE_PATH, IMAGES_FILE_PATH, OUTPUTS_FOLDER
5 from src.solutions import q2, q3, q4, q5, q6
6 from src.solutions.q2 import LinearRegressionParameters
7 from src.generate_images import generate_images
8
9
10 if __name__ == "__main__":
11     if not os.path.exists(OUTPUTS_FOLDER):
12         os.makedirs(OUTPUTS_FOLDER)
13
14     # # Question 2
15     # Q2.OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q2")
16     # if not os.path.exists(Q2.OUTPUT_FOLDER):
17     #     os.makedirs(Q2.OUTPUT_FOLDER)
18     # with open(CO2_FILE_PATH) as file:
19     #     lines = [line.rstrip().split() for line in file]
20     #
21     # df_co2 = pd.DataFrame(
22     #     np.array([line for line in lines if line[0] != "#"]).astype(float)
23     # )
24     # column_names = lines[max([i for i, line in enumerate(lines) if line[0] == "#"])[1:]]
25     # df_co2.columns = column_names
26     # t = df_co2.decimal.values[:] - np.min(df_co2.decimal.values[:])
27     # y = df_co2.average.values[:].reshape(1, -1)
28     #
29     # sigma = 1
30     # mean = np.array([0, 360]).reshape(-1, 1)
31     # covariance = np.array(
32     #     [
33     #         [10**2, 0],
34     #         [0, 100**2],
35     #     ]
36     # )
37     #
38     # prior_linear_regression_parameters = LinearRegressionParameters(
39     #     mean=mean,
40     #     covariance=covariance,
41     # )
42     # posterior_linear_regression_parameters = q2.a(
43     #     t,
44     #     y,
45     #     sigma,
46     #     prior_linear_regression_parameters,
47     #     save_path=os.path.join(Q2.OUTPUT_FOLDER, "a"),
48     # )
49     # q2.b(
50     #     t_year=df_co2.decimal.values[:],
51     #     t=t,
52     #     y=y,
53     #     linear_regression_parameters=posterior_linear_regression_parameters,
54     #     error_mean=0,
55     #     error_variance=1,
56     #     save_path=os.path.join(Q2.OUTPUT_FOLDER, "b"),
57     # )
58
59     # Question 3
60     Q3.OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q3")
61     if not os.path.exists(Q3.OUTPUT_FOLDER):
62         os.makedirs(Q3.OUTPUT_FOLDER)
63
64     q3.learn_binary_factors(
65         x=generate_images(),
66         k=8,
67         em_maximum_iterations=5,
68         e_maximum_steps=100,
69         e_convergence_criterion=0,
70     )
```

main.py