

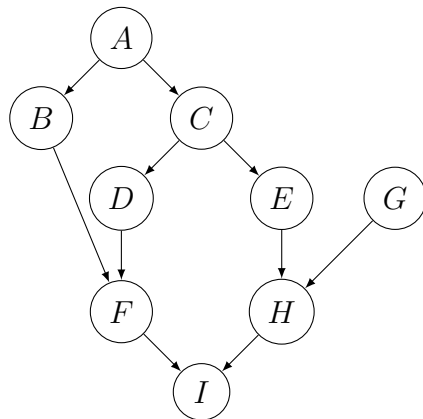
COMP0085 Summative Assignment

Jan 4, 2023

Question 1

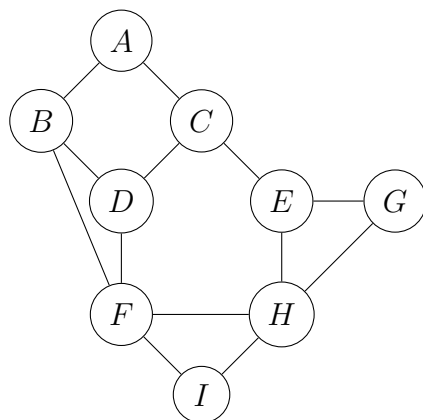
(a)

The directed acyclic graph:

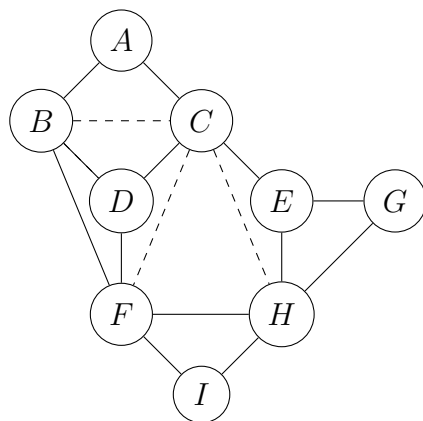


(b)

The moralised graph:

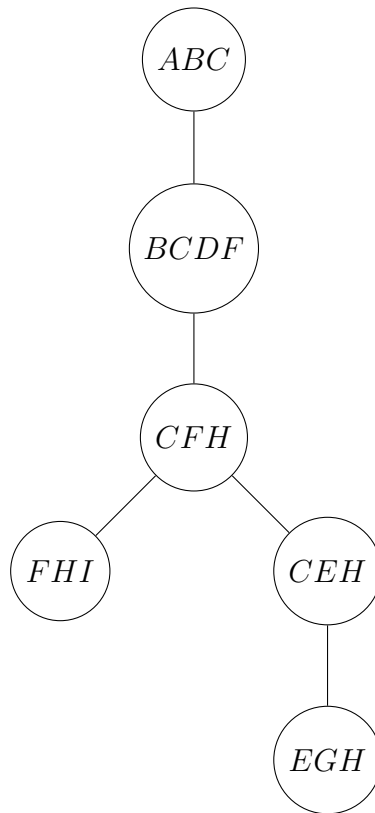


An effective triangulation:



where the dashed lines are edges added to triangulate the moralised graph.

The resulting junction tree:



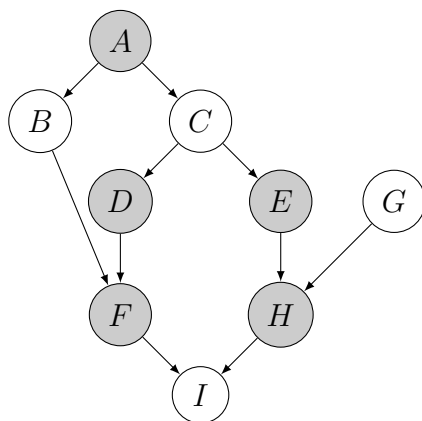
where the circular nodes are cliques.

The junction tree redrawn as a factor graph:



where the circular nodes are cliques and the square nodes are separators/factors.

(c)



The set $\{A, D, E, F, H\}$ is a non-unique smallest set of molecules such that if the concentrations of the species within the set are known, the concentrations of the others $\{B, C, G, I\}$ would all be independent (conditioned on the measured ones).

(d)

(e)

Question 2

(a)

We want the posterior mean and covariance over a and b . Defining a weight vector \mathbf{w} :

$$\mathbf{w} = \begin{bmatrix} a \\ b \end{bmatrix}$$

Our distribution for \mathbf{w} :

$$P(\mathbf{w}) = \mathcal{N} \left(\begin{bmatrix} \mu_a \\ \mu_b \end{bmatrix}, \begin{bmatrix} \sigma_a^2 & 0 \\ 0 & \sigma_b^2 \end{bmatrix} \right) = \mathcal{N}(\mu_{\mathbf{w}}, \Sigma_{\mathbf{w}})$$

Moreover, for our data $\mathcal{D} = \{\mathbf{X}, \mathbf{Y}\}$:

$$P(\mathcal{D}|\mathbf{w}) = \mathcal{N}(\mathbf{Y} - \mathbf{w}^T \mathbf{X}, \sigma^2 \mathbf{I})$$

where $\mathbf{X} = \begin{bmatrix} t_1 & t_2 \cdots t_N \\ 1 & 1 \cdots 1 \end{bmatrix} \in \mathbb{R}^{2 \times N}$ and $\mathbf{Y} \in \mathbb{R}^{1 \times N}$.

Knowing:

$$P(\mathbf{w}|\mathcal{D}) \propto P(\mathcal{D}|\mathbf{w})P(\mathbf{w})$$

we can substitute the above distributions:

$$P(\mathbf{w}|\mathcal{D}) \propto \exp \left(\frac{-1}{2\sigma^2} (\mathbf{Y} - \mathbf{w}^T \mathbf{X}) (\mathbf{Y} - \mathbf{w}^T \mathbf{X})^T \right) \exp \left(\frac{-1}{2} (\mathbf{w} - \mu_{\mathbf{w}})^T \Sigma_{\mathbf{w}}^{-1} (\mathbf{w} - \mu_{\mathbf{w}}) \right)$$

expanding:

$$\log P(\mathbf{w}|\mathcal{D}) \propto \frac{-1}{2} \left(\frac{\mathbf{Y}\mathbf{Y}^T}{\sigma^2} - 2\mathbf{w}^T \frac{\mathbf{X}\mathbf{Y}^T}{\sigma^2} + \mathbf{w}^T \frac{\mathbf{X}\mathbf{X}^T}{\sigma^2} \mathbf{w} + \mathbf{w}^T \Sigma_{\mathbf{w}}^{-1} \mathbf{w} - 2\mathbf{w}^T \Sigma_{\mathbf{w}}^{-1} \mu_{\mathbf{w}} + \mu_{\mathbf{w}}^T \Sigma_{\mathbf{w}}^{-1} \mu_{\mathbf{w}} \right)$$

collecting \mathbf{w} terms:

$$\log P(\mathbf{w}|\mathcal{D}) \propto \frac{-1}{2} \left(\mathbf{w}^T \left(\frac{\mathbf{X}\mathbf{X}^T}{\sigma^2} + \Sigma_{\mathbf{w}}^{-1} \right) \mathbf{w} - 2\mathbf{w}^T \left(\frac{\mathbf{X}\mathbf{Y}^T}{\sigma^2} + \Sigma_{\mathbf{w}}^{-1} \mu_{\mathbf{w}} \right) \right)$$

Knowing that the posterior $P(\mathbf{w}|\mathcal{D})$ will be Gaussian with mean $\bar{\mu}_w$ and covariance $\bar{\Sigma}_w$, we can see that expanding the exponent component would have the form:

$$(\mathbf{w} - \bar{\mu}_w)^T \bar{\Sigma}_w^{-1} (\mathbf{w} - \bar{\mu}_w) = \mathbf{w}^T \bar{\Sigma}_w^{-1} \mathbf{w} - 2\mathbf{w}^T \bar{\Sigma}_w^{-1} \bar{\mu}_w + \bar{\mu}_w^T \bar{\Sigma}_w^{-1} \bar{\mu}_w$$

Thus we can identify the posterior covariance:

$$\bar{\Sigma}_w = \left(\frac{\mathbf{X}\mathbf{X}^T}{\sigma^2} + \Sigma_{\mathbf{w}}^{-1} \right)^{-1}$$

and the posterior mean:

$$\bar{\mu}_w = \bar{\Sigma}_w \left(\frac{\mathbf{X}\mathbf{Y}^T}{\sigma^2} + \Sigma_{\mathbf{w}}^{-1} \mu_{\mathbf{w}} \right)$$

The Python code:

```

1 import scipy
2 import numpy as np
3 from dataclasses import dataclass
4 import pandas as pd
5 import dataframe_image as dfi
6 import matplotlib.pyplot as plt
7
8 @dataclass
9 class LinearRegressionParameters:
10     mean: np.ndarray
11     covariance: np.ndarray
12
13     @property
14     def precision(self):
15         return np.linalg.inv(self.covariance)
16
17     def predict(self, x: np.ndarray) -> np.ndarray:
18         return self.mean.T @ x
19
20
21 @dataclass
22 class Theta:
23     linear_regression_parameters: LinearRegressionParameters
24     sigma: float
25
26     @property
27     def variance(self):
28         return self.sigma**2
29
30     @property
31     def precision(self):
32         return 1 / self.variance
33
34
35 def compute_posterior(
36     x: np.ndarray,
37     y: np.ndarray,
38     prior_linear_regression_parameters: LinearRegressionParameters,
39     residuals_precision: float,
40 ) -> LinearRegressionParameters:
41     """
42     Compute the parameters of the posterior distribution on the linear regression weights
43
44     :param x: design matrix (number of features, number of data points)
45     :param y: response matrix (1, number of data points)
46     :param prior_linear_regression_parameters: parameters for the prior distribution on the linear regression
47           weights
48     :param residuals_precision: the precision of the residuals of the linear regression
49     :return: parameters for the posterior distribution on the linear regression weights
50     """
51     posterior_covariance = np.linalg.inv(
52         residuals_precision * x @ x.T + prior_linear_regression_parameters.precision
53     )
54     posterior_mean = posterior_covariance @ (
55         residuals_precision * x @ y.T
56         + prior_linear_regression_parameters.precision
57         @ prior_linear_regression_parameters.mean
58     )
59     return LinearRegressionParameters(
60         mean=posterior_mean, covariance=posterior_covariance
61     )
62
63 def construct_design_matrix(t: np.ndarray):
64     return np.stack((t, np.ones(t.shape)), axis=1).T
65
66
67 def a(
68     t: np.ndarray,
69     y: np.ndarray,
70     sigma: float,
71     prior_linear_regression_parameters: LinearRegressionParameters,
72     save_path: str
73 ) -> LinearRegressionParameters:
74     x = construct_design_matrix(t)
75     prior_theta = Theta(
76         linear_regression_parameters=prior_linear_regression_parameters,
77         sigma=sigma,
78     )
79     posterior_linear_regression_parameters = compute_posterior(
80         x,
81         y,
82         prior_linear_regression_parameters,
83         residuals_precision=prior_theta.precision,
84     )
85     df_mean = pd.DataFrame(posterior_linear_regression_parameters.mean, columns=["value"])
86     df_mean.index = ["a", "b"]
87     df_mean = pd.concat([df_mean], keys=["parameters"])
88     dfi.export(df_mean, save_path + "-mean.png")
89
90     df_covariance = pd.DataFrame(posterior_linear_regression_parameters.covariance, columns=["a", "b"])
91     df_covariance.index = ["a", "b"]
92     df_covariance = pd.concat([df_covariance], keys=["parameters"])
93     df_covariance = pd.concat([df_covariance.T], keys=["parameters"])

```



```

94     dfi.export(df_covariance, save_path + "-covariance.png")
95     return posterior_linear_regression_parameters
96
97
98 def b(t_year, t, y, linear_regression_parameters, error_mean, error_variance, save_path):
99     x = construct_design_matrix(t)
100     residuals = y - linear_regression_parameters.predict(x)
101     plt.plot(t_year.reshape(-1), residuals.reshape(-1))
102     plt.xlabel("date (decimal year)")
103     plt.ylabel("residual")
104     plt.title("2b: g_obs(t)")
105     plt.savefig(save_path + "-residuals-timeseries")
106     plt.close()
107
108     count, bins = np.histogram(residuals, bins=100, density=True)
109     plt.bar(bins[1:], count, label="residuals")
110     plt.plot(bins[1:], scipy.stats.norm.pdf(bins[1:], loc=error_mean, scale=error_variance), color="red", label="
111             e(t)")
111     plt.xlabel("residual bin")
112     plt.ylabel("density")
113     plt.title("2b: Residuals Density")
114     plt.legend()
115     plt.savefig(save_path + "-residuals-density-estimation")
116     plt.close()

```

src/solutions/q2.py

Question 3

(a)

The free energy is can be calculated as:

$$\mathcal{F}(q, \theta) = \langle \log P(\mathbf{x}, \mathbf{s} | \theta) \rangle_{q(\mathbf{s})} + H[Q(\mathbf{s})]$$

Knowing,

$$\log P(\mathbf{x}, \mathbf{s} | \theta) = \log P(\mathbf{x} | \mathbf{s}, \theta) + \log P(\mathbf{s} | \theta)$$

we can write:

$$\mathcal{F}(Q, \theta) = \langle \log P(\mathbf{x} | \mathbf{s}, \theta) \rangle_{q(\mathbf{s})} + \langle \log P(\mathbf{s} | \theta) \rangle_{q(\mathbf{s})} + H[q(\mathbf{s})]$$

Moreover, our mean field approximation:

$$q(\mathbf{s}) = \prod_{i=1}^K q_i(s_i)$$

where $q_i(s_i) = \lambda_i^{s_i} (1 - \lambda_i)^{(1-s_i)}$.

To compute the first term:

$$P(\mathbf{x} | \mathbf{s}, \theta) = \mathcal{N} \left(\sum_{i=1}^K s_i \mu_i, \sigma^2 \mathbf{I} \right)$$

substituting the appropriate terms:

$$P(\mathbf{x} | \mathbf{s}, \theta) = 2\pi^{-\frac{d}{2}} |\sigma^2 \mathbf{I}|^{-\frac{1}{2}} \exp \left(-\frac{1}{2} \left(\mathbf{x} - \sum_{i=1}^K s_i \mu_i \right)^T \frac{1}{\sigma^2} \mathbf{I} \left(\mathbf{x} - \sum_{i=1}^K s_i \mu_i \right) \right)$$

with d being the number of dimensions.

Taking the logarithm:

$$\log P(\mathbf{x} | \mathbf{s}, \theta) = -\frac{d}{2} \log(2\pi) - \log(\sigma) - \frac{1}{2\sigma^2} \left(\mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K s_i \mu_i + \sum_{i=1}^K \sum_{j=1}^K s_i s_j \mu_i^T \mu_j \right)$$

The expectation distributed to the relevant terms:

$$\langle \log P(\mathbf{x} | \mathbf{s}, \theta) \rangle_{q(\mathbf{s})} = -\frac{d}{2} \log(2\pi) - \log(\sigma) - \frac{1}{2\sigma^2} \left(\mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K \langle s_i \rangle_{q_i(s_i)} \mu_i + \sum_{i=1}^K \sum_{j=1}^K \langle s_i s_j \rangle_{q_i(s_i) q_j(s_j)} \mu_i^T \mu_j \right)$$

Evaluating the expectations:

$$\langle \log P(\mathbf{x} | \mathbf{s}, \theta) \rangle_{q(\mathbf{s})} = -\frac{d}{2} \log(2\pi) - \log(\sigma) - \frac{1}{2\sigma^2} \left(\mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K \lambda_i \mu_i + \sum_{i=1}^K \sum_{j=1, j \neq i}^K \lambda_i \lambda_j \mu_i^T \mu_j + \sum_{i=1}^K \lambda_i \mu_i^T \mu_i \right)$$

where $\langle s_i s_i \rangle_{q_i(s_i)} = \langle s_i \rangle_{q_i(s_i)}$ because $s_i \in \{0, 1\}$.
To compute the second term:

$$P(\mathbf{s}|\theta) = \prod_{i=1}^K \pi_i^{s_i} (1 - \pi_i)^{(1-s_i)}$$

Taking the logarithm:

$$\log P(\mathbf{s}|\theta) = \sum_{i=1}^K s_i \log \pi_i + (1 - s_i) \log(1 - \pi_i)$$

The expectation distributed to the relevant terms:

$$\langle \log P(\mathbf{s}|\theta) \rangle_{q(\mathbf{s})} = \sum_{i=1}^K \langle s_i \rangle_{q_i(s_i)} \log \pi_i + (1 - \langle s_i \rangle_{q_i(s_i)}) \log(1 - \pi_i)$$

Evaluating the expectations:

$$\langle \log P(\mathbf{s}|\theta) \rangle_{q(\mathbf{s})} = \sum_{i=1}^K \lambda_i \log \pi_i + (1 - \lambda_i) \log(1 - \pi_i)$$

To compute the third term, we use the mean field factorisation:

$$H[q(\mathbf{s})] = \sum_{i=1}^K H[q_i(s_i)]$$

Thus,

$$H[q(\mathbf{s})] = - \sum_{i=1}^K \sum_{s_i \in \{0,1\}} q_i(s_i) \log q_i(s_i)$$

Substituting the appropriate values:

$$H[q(\mathbf{s})] = - \sum_{i=1}^K \lambda_i \log \lambda_i + (1 - \lambda_i) \log(1 - \lambda_i)$$

Combining, we have our free energy expression:

$$\begin{aligned} \mathcal{F}(q, \theta) = & \\ & \frac{-d}{2} \log(2\pi) - \log(\sigma) - \frac{1}{2\sigma^2} \left(\mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K \lambda_i \mu_i + \sum_{i=1}^K \sum_{j=1, j \neq i}^K \lambda_i \lambda_j \mu_i^T \mu_j + \sum_{i=1}^K \lambda_i \mu_i^T \mu_i \right) \\ & + \sum_{i=1}^K \lambda_i \log \pi_i + (1 - \lambda_i) \log(1 - \pi_i) \\ & - \sum_{i=1}^K \lambda_i \log \lambda_i + (1 - \lambda_i) \log(1 - \lambda_i) \end{aligned}$$

To derive the partial update for $q_i(s_i)$ we take the variational derivative of the Lagrangian, enforcing the normalisation of q_i :

$$\frac{\partial}{\partial q_i} \left(\mathcal{F}(q, \theta) + \lambda^{LG} \int q_i - 1 \right) = \langle \log P(\mathbf{x}, \mathbf{s}|\theta) \rangle_{\prod_{j \neq i} q_j(s_j)} - \log q_i(s_i) - 1 + \lambda^{LG}$$

where λ^{LG} is the Lagrange multiplier.

Setting this to zero we can solve for the λ_i that maximises the free energy:

$$\log q_i(s_i) = \langle \log P(\mathbf{x}, \mathbf{s}|\theta) \rangle_{\prod_{j \neq i} q_j(s_j)} - 1 + \lambda^{LG}$$

Similar to our free energy derivation:

$$\langle \log P(\mathbf{x}|\mathbf{s}, \theta) \rangle_{\prod_{j \neq i} q_j(s_j)} \propto -\frac{1}{2\sigma^2} \left(\mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{k=1}^K \langle s_k \rangle_{\prod_{j \neq i} q_j(s_j)} \mu_k + \sum_{k=1}^K \sum_{j=1}^K \langle s_k s_j \rangle_{\prod_{j \neq i} q_j(s_j)} \right)$$

and

$$\langle \log P(\mathbf{s}|\theta) \rangle_{\prod_{j \neq i} q_j(s_j)} = \sum_{k=1}^K \langle s_k \rangle_{\prod_{j \neq i} q_j(s_j)} \log \pi_k + (1 - \langle s_k \rangle_{\prod_{j \neq i} q_j(s_j)}) \log(1 - \pi_k)$$

We can write:

$$\log q_i(s_i) \propto \log P(\mathbf{x}|\mathbf{s}, \theta)_{\prod_{j \neq i} q_j(s_j)} + \langle \log P(\mathbf{s}|\theta) \rangle_{\prod_{j \neq i} q_j(s_j)}$$

Substituting the relevant terms:

$$\log q_i(s_i) \propto -\frac{1}{2\sigma^2} \left(-2s_i \mathbf{x}^T \mu_i + s_i s_i \mu_i^T \mu_i + 2 \sum_{j=1, j \neq i}^K s_i \lambda_j \mu_i^T \mu_j \right) + s_i \log \pi_i + (1 - s_i) \log(1 - \pi_i)$$

Knowing $\log q_i(s_i) = s_i \log \lambda_i + (1 - s_i) \log(1 - \lambda_i)$:

$$\log q_i(s_i) \propto s_i \log \frac{\lambda_i}{1 - \lambda_i}$$

Thus,

$$s_i \log \frac{\lambda_i}{1 - \lambda_i} \propto -\frac{1}{2\sigma^2} \left(-2s_i \mathbf{x}^T \mu_i + s_i s_i \mu_i^T \mu_i + 2 \sum_{j=1, j \neq i}^K s_i \lambda_j \mu_i^T \mu_j \right) + s_i \log \frac{\pi_i}{1 - \pi_i}$$

Also, because $s_i \in \{0, 1\}$ we know that $s_i^2 = s_i$:

$$s_i \log \frac{\lambda_i}{1 - \lambda_i} \propto -\frac{1}{2\sigma^2} \left(-2s_i \mathbf{x}^T \mu_i + s_i \mu_i^T \mu_i + 2 \sum_{j=1, j \neq i}^K s_i \lambda_j \mu_i^T \mu_j \right) + s_i \log \frac{\pi_i}{1 - \pi_i}$$

Because we have only kept terms with s_i , this is an equality:

$$s_i \log \frac{\lambda_i}{1 - \lambda_i} = \frac{s_i \mu_i^T}{2\sigma^2} \left(2\mathbf{x} - \mu_i - 2 \sum_{j=1, j \neq i}^K \lambda_j \mu_j \right) + s_i \log \frac{\pi_i}{1 - \pi_i}$$

Solving for λ_i :

$$\lambda_i = \frac{1}{1 + \exp \left[- \left(\frac{\mu_i^T}{\sigma^2} \left(\mathbf{x} - \frac{\mu_i}{2} - \sum_{j=1, j \neq i}^K \lambda_j \mu_j \right) + \log \frac{\pi_i}{1 - \pi_i} \right) \right]}$$

we have our partial update.

The Python code:

```

1 from dataclasses import dataclass
2
3 import numpy as np
4
5 from demo_code.MStep import m_step
6
7
8 @dataclass
9 class MeanFieldApproximation:
10     """
11     lambda_matrix: parameters variational approximation (number_of_points, number_of_latent_variables)
12     """
13
14     lambda_matrix: np.ndarray
15
16     def lambda_matrix_exclude(self, exclude_latent_index: int) -> np.ndarray:
17         # (number_of_points, number_of_latent_variables-1)
18         return np.concatenate(
19             (
20                 self.lambda_matrix[:, :exclude_latent_index],
21                 self.lambda_matrix[:, exclude_latent_index + 1 :],
22             ),
23             axis=1,
24         )
25
26     @property
27     def log_lambda_matrix(self):
28         return np.log(self.lambda_matrix)
29
30     @property
31     def log_one_minus_lambda_matrix(self):
32         return np.log(1 - self.lambda_matrix)
33
34     @property
35     def n(self):
36         return self.lambda_matrix.shape[0]
37
38     @property
39     def k(self):
40         return self.lambda_matrix.shape[1]
41
42
43 def init_mean_field_approximation(k: int, n: int) -> MeanFieldApproximation:
44     return MeanFieldApproximation(
45         lambda_matrix=np.random.random(size=(n, k)),
46     )
47
48
49 @dataclass
50 class BinaryLatentFactorModel:
51     """
52     mu: matrix of means (number_of_dimensions, number_of_latent_variables)
53     sigma: gaussian noise parameter
54     pi: vector of priors (1, number_of_latent_variables)
55     """
56
57     mu: np.ndarray
58     sigma: float
59     pi: np.ndarray
60
61     def mu_exclude(self, exclude_latent_index: int) -> np.ndarray:
62         # (number_of_dimensions, number_of_latent_variables-1)
63         return np.concatenate(
64             (self.mu[:, :exclude_latent_index], self.mu[:, exclude_latent_index + 1 :]),
65             axis=1,
66         )
67
68     @property
69     def log_pi(self):
70         return np.log(self.pi)
71
72     @property
73     def log_one_minus_pi(self):
74         return np.log(1 - self.pi)
75
76     @property
77     def variance(self):
78         return self.sigma**2
79
80     @property
81     def precision(self):
82         return 1 / self.variance
83
84     @property
85     def d(self):
86         return self.mu.shape[0]
87
88     @property
89     def k(self):
90         return self.mu.shape[1]
91
92
93 def init_binary_latent_factor_model(
94     x: np.ndarray,

```

```

95     mean_field_approximation: MeanFieldApproximation,
96 ) -> BinaryLatentFactorModel:
97     return maximisation_step(x, mean_field_approximation)
98
99
100 def _compute_expectation_log_p_x_s_given_theta(
101     x: np.ndarray,
102     binary_latent_factor_model: BinaryLatentFactorModel,
103     mean_field_approximation: MeanFieldApproximation,
104 ) -> float:
105     """
106     The first term of the free energy, the expectation of log P(X,S|theta)
107
108     :param x: data matrix (number_of_points, number_of_dimensions)
109     :param binary_latent_factor_model: a binary_latent_factor_model
110     :param mean_field_approximation: a mean_field_approximation
111     :return: the expectation of log P(X,S|theta)
112     """
113     # (number_of_points, number_of_dimensions)
114     mu_lambda = mean_field_approximation.lambda_matrix @ binary_latent_factor_model.mu.T
115
116     # (number_of_latent_variables, number_of_latent_variables)
117     expectation_s_i_s_j_mu_i_mu_j = np.multiply(
118         mean_field_approximation.lambda_matrix.T
119         @ mean_field_approximation.lambda_matrix,
120         binary_latent_factor_model.mu.T @ binary_latent_factor_model.mu,
121     )
122
123     expectation_log_p_x_given_s_theta = (
124         (-binary_latent_factor_model.d / 2) * np.log(2 * np.pi)
125         - np.log(binary_latent_factor_model.sigma)
126         - (0.5 * binary_latent_factor_model.precision)
127         * (
128             np.sum(np.multiply(x, x))
129             - 2 * np.sum(np.multiply(x, mu_lambda))
130             + np.sum(expectation_s_i_s_j_mu_i_mu_j)
131             - np.trace(
132                 expectation_s_i_s_j_mu_i_mu_j
133             ) # remove incorrect E[s_i s_i] = lambda_i * lambda_i
134             + np.sum( # add correct E[s_i s_i] = lambda_i
135                 mean_field_approximation.lambda_matrix
136                 @ np.multiply(
137                     binary_latent_factor_model.mu, binary_latent_factor_model.mu
138                 ).T
139             )
140         )
141     )
142     expectation_log_p_s_given_theta = np.sum(
143         mean_field_approximation.lambda_matrix @ binary_latent_factor_model.log_pi.T
144         + (1 - mean_field_approximation.lambda_matrix)
145         @ binary_latent_factor_model.log_one_minus_pi.T
146     )
147     return expectation_log_p_x_given_s_theta + expectation_log_p_s_given_theta
148
149
150 def _compute_mean_field_approximation_entropy(
151     mean_field_approximation: MeanFieldApproximation,
152 ) -> float:
153     return -np.sum(
154         np.multiply(
155             mean_field_approximation.lambda_matrix,
156             mean_field_approximation.log_lambda_matrix,
157         )
158         + np.multiply(
159             1 - mean_field_approximation.lambda_matrix,
160             mean_field_approximation.log_one_minus_lambda_matrix,
161         )
162     )
163
164
165 def compute_free_energy(
166     x: np.ndarray,
167     binary_latent_factor_model: BinaryLatentFactorModel,
168     mean_field_approximation: MeanFieldApproximation,
169 ) -> float:
170     """
171     free energy associated with current EM parameters and data x
172
173     :param x: data matrix (number_of_points, number_of_dimensions)
174     :param binary_latent_factor_model: a binary_latent_factor_model
175     :param mean_field_approximation: a mean_field_approximation
176     :return: free energy
177     """
178     expectation_log_p_x_s_given_theta = _compute_expectation_log_p_x_s_given_theta(
179         x, binary_latent_factor_model, mean_field_approximation
180     )
181     mean_field_approximation_entropy = _compute_mean_field_approximation_entropy(
182         mean_field_approximation
183     )
184     return expectation_log_p_x_s_given_theta + mean_field_approximation_entropy
185
186
187 def partial_expectation_step(
188     x: np.ndarray,
189     binary_latent_factor_model: BinaryLatentFactorModel,
190     mean_field_approximation: MeanFieldApproximation,

```

```

191 latent_factor: int,
192 ) -> np.ndarray:
193     """ Partial Variational E step for factor i for all data points
194
195     :param x: data matrix (number_of_points, number_of_dimensions)
196     :param binary_latent_factor_model: a binary_latent_factor_model
197     :param mean_field_approximation: a mean_field_approximation
198     :param latent_factor: latent factor to compute partial update
199     :return: lambda_vector: new lambda parameters for the latent factor (number_of_points, 1)
200     """
201     lambda_matrix_excluded = mean_field_approximation.lambda_matrix_exclude(
202         latent_factor
203     )
204     mu_excluded = binary_latent_factor_model.mu_exclude(latent_factor)
205
206     mu_latent = binary_latent_factor_model.mu[:, latent_factor]
207     # (number_of_points, 1)
208     partial_expectation_log_p_x_given_s_theta_proportion = (
209         binary_latent_factor_model.precision
210         * (
211             x # (number_of_points, number_of_dimensions)
212             - 0.5 * mu_latent.T # (1, number_of_dimensions)
213             - lambda_matrix_excluded # (number_of_points, number_of_latent_variables-1)
214             @ mu_excluded.T # (number_of_latent_variables-1, number_of_dimensions)
215         )
216         @ mu_latent # (number_of_dimensions, 1)
217     )
218
219     # (1, 1)
220     partial_expectation_log_p_s_given_theta_proportion = np.log(
221         binary_latent_factor_model.pi[0, latent_factor]
222         / (1 - binary_latent_factor_model.pi[0, latent_factor])
223     )
224
225     # (number_of_points, 1)
226     partial_expectation_log_p_x_s_given_theta_proportion = (
227         partial_expectation_log_p_x_given_s_theta_proportion
228         + partial_expectation_log_p_s_given_theta_proportion
229     )
230
231     # (number_of_points, 1)
232     lambda_vector = 1 / (
233         1 + np.exp(-partial_expectation_log_p_x_s_given_theta_proportion)
234     )
235     lambda_vector[lambda_vector == 0] = 1e-10
236     lambda_vector[lambda_vector == 1] = 1 - 1e-10
237     return lambda_vector
238
239
240 def variational_expectation_step(
241     x: np.ndarray,
242     binary_latent_factor_model: BinaryLatentFactorModel,
243     mean_field_approximation: MeanFieldApproximation,
244     max_steps: int,
245     convergence_criterion: float,
246 ) -> MeanFieldApproximation:
247     """ Variational E step
248
249     :param x: data matrix (number_of_points, number_of_dimensions)
250     :param binary_latent_factor_model: a binary_latent_factor_model
251     :param mean_field_approximation: a mean_field_approximation
252     :param max_steps: maximum number of steps of fixed point equations
253     :param convergence_criterion: early stopping if change in free energy < convergence_criterion
254     :return: mean field approximation
255     """
256     previous_free_energy = compute_free_energy(
257         x, binary_latent_factor_model, mean_field_approximation
258     )
259     for i in range(max_steps):
260         for latent_factor in range(binary_latent_factor_model.k):
261             mean_field_approximation.lambda_matrix[
262                 :, latent_factor
263             ] = partial_expectation_step(
264                 x, binary_latent_factor_model, mean_field_approximation, latent_factor
265             )
266         free_energy = compute_free_energy(
267             x, binary_latent_factor_model, mean_field_approximation
268         )
269         if free_energy - previous_free_energy <= convergence_criterion:
270             break
271         previous_free_energy = free_energy
272     return mean_field_approximation
273
274
275 def maximisation_step(
276     x: np.ndarray,
277     mean_field_approximation: MeanFieldApproximation,
278 ) -> BinaryLatentFactorModel:
279     expectation_s = mean_field_approximation.lambda_matrix
280     expectation_ss = (
281         mean_field_approximation.lambda_matrix.T
282         @ mean_field_approximation.lambda_matrix
283     )
284     np.fill_diagonal(expectation_ss, mean_field_approximation.lambda_matrix.sum(axis=0))
285     mu, sigma, pi = m_step(x, expectation_s, expectation_ss)
286     return BinaryLatentFactorModel(

```

```

287         mu=mu,
288         sigma=sigma,
289         pi=pi,
290     )
291
292
293 def learn_binary_factors(
294     x: np.ndarray,
295     k: int,
296     em_maximum_iterations: int,
297     e_maximum_steps: int,
298     e_convergence_criterion: float,
299 ):
300     n = x.shape[0]
301     mean_field_approximation = init_mean_field_approximation(k, n)
302     binary_latent_factor_model = init_binary_latent_factor_model(
303         x, mean_field_approximation
304     )
305
306     for _ in range(em_maximum_iterations):
307         mean_field_approximation = variational_expectation_step(
308             x=x,
309             binary_latent_factor_model=binary_latent_factor_model,
310             mean_field_approximation=mean_field_approximation,
311             max_steps=e_maximum_steps,
312             convergence_criterion=e_convergence_criterion,
313         )
314         binary_latent_factor_model = maximisation_step(
315             x=x,
316             mean_field_approximation=mean_field_approximation,
317         )
318     return mean_field_approximation, binary_latent_factor_model

```

src/solutions/q3.py

Question 4

Question 5

(a)

The log-joint probability for a single observation-source pair:

$$\log p(\mathbf{s}, \mathbf{x}) = \log p(\mathbf{s}) + (\mathbf{x}|\mathbf{s})$$

Knowing $p(\mathbf{s}) = \prod_{i=1}^K p(s_i|\pi_i)$ and $p(\mathbf{x}|\mathbf{s}) = \mathcal{N}(\sum_{i=1}^K s_i \mu_i, \sigma^2 \mathbf{I})$:

$$\log p(\mathbf{s}, \mathbf{x}) \propto \frac{-1}{2} \left(\mathbf{x} - \sum_{i=1}^K s_i \mu_i \right)^T \frac{1}{\sigma^2} \mathbf{I} \left(\mathbf{x} - \sum_{i=1}^K s_i \mu_i \right) + \sum_{i=1}^K (s_i \log \pi_i + (1 - s_i) \log(1 - \pi_i))$$

Expanding:

$$\log p(\mathbf{s}, \mathbf{x}) \propto \frac{-1}{2\sigma^2} \left(\mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T \sum_{i=1}^K s_i \mu_i + \sum_{i=1}^K \sum_{j=1}^K s_i s_j \mu_i^T \mu_j \right) + \sum_{i=1}^K (s_i \log \pi_i + (1 - s_i) \log(1 - \pi_i))$$

Collecting terms pertaining to s_i :

$$\log p(\mathbf{s}, \mathbf{x}) = \sum_{i=1}^K \left(\left(\frac{\mathbf{x}^T \mu_i}{\sigma^2} + \log \frac{\pi_i}{1 - \pi_i} \right) s_i \right) + \sum_{i=1}^K \sum_{j=1}^K \left(\frac{-\mu_i^T \mu_j}{2\sigma^2} s_i s_j \right) + C$$

where C are all other terms without s_i .

Knowing that $s_i^2 = s_i$:

$$\log p(\mathbf{s}, \mathbf{x}) = \sum_{i=1}^K \left(\left(\frac{\mathbf{x}^T \mu_i}{\sigma^2} + \log \frac{\pi_i}{1 - \pi_i} - \frac{\mu_i^T \mu_j}{2\sigma^2} \right) s_i \right) + \sum_{i=1}^K \sum_{j=1}^{i-1} \left(\frac{-\mu_i^T \mu_j}{\sigma^2} s_i s_j \right) + C$$

Thus:

$$\log p(\mathbf{s}, \mathbf{x}) = \sum_{i=1}^K \log f_i(s_i) + \sum_{i=1}^K \sum_{j=1}^{i-1} \log g_{ij}(s_i, s_j)$$

where the factors are defined:

$$\log f_i(s_i) = \left(\frac{\mathbf{x}^T \mu_i}{\sigma^2} + \log \frac{\pi_i}{1 - \pi_i} - \frac{\mu_i^T \mu_j}{2\sigma^2} \right) s_i$$

and

$$\log g_{ij}(s_i, s_j) = \frac{-\mu_i^T \mu_j}{\sigma^2} s_i s_j$$

as required. Note that C can simply be absorbed into any one of these factors.

The Boltzmann Machine can be defined:

$$P(\mathbf{s}|\mathbf{W}, \mathbf{b}) = \frac{1}{Z} \exp \left(\sum_{i=1}^K \sum_{j=1}^{i-1} W_{ij} s_i s_j - \sum_{i=1}^K b_i s_i \right)$$

where $s_i \in \{0, 1\}$, the same as our source variables.

From our factorisation, we can see that $p(\mathbf{s}, \mathbf{x})$ is a Boltzmann Machine with:

$$W_{ij} = \frac{-\mu_i^T \mu_j}{\sigma^2}$$

and

$$b_i = - \left(\frac{\mathbf{x}^T \mu_i}{\sigma^2} + \log \frac{\pi_i}{1 - \pi_i} - \frac{\mu_i^T \mu_j}{2\sigma^2} \right)$$

and

$$\log Z = -C$$

(b)

For $f_i(s_i)$, we will choose a Bernoulli approximation:

$$\tilde{f}_i(s_i) = \lambda_i^{s_i} + (1 - \lambda_i)^{1-s_i}$$

Thus,

$$\log \tilde{f}_i(s_i) \propto \log \left(\frac{\lambda_i}{1 - \lambda_i} \right) s_i$$

For $g_{ij}(s_i, s_j)$, we will choose a product of Bernoulli's approximation:

$$\tilde{g}_{ij}(s_i, s_j) = (\theta_i^{s_i} + (1 - \theta_i)^{1-s_i}) (\theta_j^{s_j} + (1 - \theta_j)^{1-s_j})$$

Thus,

$$\log \tilde{g}_{ij}(s_i, s_j) \propto \log \left(\frac{\theta_i}{1 - \theta_i} \right) s_i + \log \left(\frac{\theta_j}{1 - \theta_j} \right) s_j$$

To derive the a message passing scheme, we first define:

$$q(\mathbf{s}) = \left(\prod_{i=1}^K \tilde{f}_i(s_i) \right) \left(\prod_{i=1}^K \prod_{j=1}^{i-1} \tilde{g}_{ij}(s_i, s_j) \right)$$

Thus, we can derive cavity distributions:

$$q_{-\tilde{f}_i(s_i)}(\mathbf{s}) = \left(\prod_{i'=1, i' \neq i}^K \tilde{f}_{i'}(s_{i'}) \right) \left(\prod_{i=1}^K \prod_{j=1}^{i-1} \tilde{g}_{ij}(s_i, s_j) \right)$$

and

$$q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) = \left(\prod_{i'=1}^K \tilde{f}_{i'}(s_{i'}) \right) \left(\prod_{i'=1}^K \prod_{\substack{j'=1 \\ i', j' \neq i, j}}^{i'-1} \tilde{g}_{i'j'}(s_{i'}, s_{j'}) \right)$$

For $\tilde{f}_i(s_i)$, we do not need to make an approximation step. This is because we are minimising:

$$\tilde{f}_i(s_i) = \arg \min \mathbf{KL} \left[f_i(s_i) q_{-\tilde{f}_i(s_i)}(\mathbf{s}) \| \tilde{f}_i(s_i) q_{-\tilde{f}_i(s_i)}(\mathbf{s}) \right]$$

We know that the factor $\log f_i(s_i)$ is a Bernoulli of the form $b_i s_i$. Because our approximation is also Bernoulli, we can simply solve for λ_i in $\log \tilde{f}_i(s_i)$:

$$\log \tilde{f}_i(s_i) = \log f_i(s_i)$$

$$\log \left(\frac{\lambda_i}{1 - \lambda_i} \right) s_i = b_i s_i$$

$$\lambda_i = \frac{1}{1 + \exp(-b_i)}$$

On the other hand, for $\tilde{g}_{ij}(s_i, s_j)$, we will approximate with:

$$\tilde{g}_{ij}(s_i, s_j) = \arg \min \mathbf{KL} \left[g_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) \| \tilde{g}_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) \right]$$

Note that because $\tilde{g}_{ij}(s_i, s_j)$ is the product of two Bernoulli distributions, we only require the natural parameters:

$$\phi_{ij}(\theta) = \begin{bmatrix} \theta_i \\ \theta_j \end{bmatrix}$$

the mean with respect to s_i and s_j respectively.

We can write:

$$\begin{aligned} \log \tilde{g}_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) &\propto \log \left(\frac{\theta_i}{1 - \theta_i} \right) s_i + \log \left(\frac{\theta_j}{1 - \theta_j} \right) s_j \\ &\quad + \sum_{i'=1}^K \log \left(\frac{\lambda_{i'}}{1 - \lambda_{i'}} \right) s_{i'} \\ &\quad + \sum_{i'=1}^K \sum_{\substack{j'=1 \\ j' \neq i, j}}^{i'-1} \log \left(\frac{\theta_{i'}}{1 - \theta_{i'}} \right) s_{i'} + \log \left(\frac{\theta_{j'}}{1 - \theta_{j'}} \right) s_{j'} \end{aligned}$$

Only including terms with s_i and s_j :

$$\begin{aligned} \log \tilde{g}_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) &\propto \left((K - 1) \log \left(\frac{\theta_i}{1 - \theta_i} \right) + \log \left(\frac{\lambda_i}{1 - \lambda_i} \right) \right) s_i \\ &\quad + \left((K - 1) \log \left(\frac{\theta_j}{1 - \theta_j} \right) + \log \left(\frac{\lambda_j}{1 - \lambda_j} \right) \right) s_j \end{aligned}$$

Moreover:

$$\begin{aligned} \log g_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) &\propto \frac{-\mu_i^T \mu_j}{\sigma^2} s_i s_j \\ &\quad + \sum_{i'=1}^K \log \left(\frac{\lambda_{i'}}{1 - \lambda_{i'}} \right) s_{i'} \\ &\quad + \sum_{i'=1}^K \sum_{\substack{j'=1 \\ j' \neq i, j}}^{i'-1} \log \left(\frac{\theta_{i'}}{1 - \theta_{i'}} \right) s_{i'} + \log \left(\frac{\theta_{j'}}{1 - \theta_{j'}} \right) s_{j'} \end{aligned}$$

Only including terms with s_i and s_j :

$$\begin{aligned} \log g_{ij}(s_i, s_j) q_{-\tilde{g}_{ij}(s_i, s_j)}(\mathbf{s}) &\propto \frac{-\mu_i^T \mu_j}{\sigma^2} s_i s_j \\ &\quad + \left((K - 2) \log \left(\frac{\theta_i}{1 - \theta_i} \right) + \log \left(\frac{\lambda_i}{1 - \lambda_i} \right) \right) s_i \\ &\quad + \left((K - 2) \log \left(\frac{\theta_j}{1 - \theta_j} \right) + \log \left(\frac{\lambda_j}{1 - \lambda_j} \right) \right) s_j \end{aligned}$$

Appendix 1: constants.py

```
1 import os
2
3 DATA_FOLDER = "data"
4
5 CO2_FILE_PATH = os.path.join(DATA_FOLDER, "co2.txt")
6 IMAGES_FILE_PATH = os.path.join(DATA_FOLDER, "images.jpg")
7
8 OUTPUTS_FOLDER = "outputs"
9
10 DEFAULT_SEED = 0
11
12 M1 = [0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0]
13
14 M2 = [0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0]
15
16 M3 = [1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
17
18 M4 = [1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1]
19
20 M5 = [0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0]
21
22 M6 = [1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1]
23
24 M7 = [0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0]
25
26 M8 = [0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1]
```

src/constants.py

Appendix 2: main.py

```
1 import os
2 import pandas as pd
3 import numpy as np
4 from src.constants import CO2_FILE_PATH, IMAGES_FILE_PATH, OUTPUTS_FOLDER
5 from src.solutions import q2, q3, q4, q5, q6
6 from src.solutions.q2 import LinearRegressionParameters
7 from src.generate_images import generate_images
8
9
10 if __name__ == "__main__":
11     if not os.path.exists(OUTPUTS_FOLDER):
12         os.makedirs(OUTPUTS_FOLDER)
13
14     with open(CO2_FILE_PATH) as file:
15         lines = [line.rstrip().split() for line in file]
16
17     # Question 2
18     Q2_OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q2")
19     if not os.path.exists(Q2_OUTPUT_FOLDER):
20         os.makedirs(Q2_OUTPUT_FOLDER)
21     df_co2 = pd.DataFrame(np.array([line for line in lines if line[0] != "#"]).astype(float))
22     column_names = lines[max([i for i, line in enumerate(lines) if line[0] == "#"])[1:]]
23     df_co2.columns = column_names
24     t = df_co2.decimal.values[:] - np.min(df_co2.decimal.values[:])
25     y = df_co2.average.values[:].reshape(1, -1)
26
27     sigma = 1
28     mean = np.array([0, 360]).reshape(-1, 1)
29     covariance = np.array(
30         [
31             [10 ** 2, 0],
32             [0, 100 ** 2],
33         ]
34     )
35
36     prior_linear_regression_parameters = LinearRegressionParameters(
37         mean=mean,
38         covariance=covariance,
39     )
40     posterior_linear_regression_parameters = q2.a(t, y, sigma, prior_linear_regression_parameters, save_path=os.
41         path.join(Q2_OUTPUT_FOLDER, "a"))
42     q2.b(
43         t_year=df_co2.decimal.values[:], t=t, y=y, linear_regression_parameters=
44         posterior_linear_regression_parameters,
45         error_mean = 0, error_variance=1, save_path= os.path.join(Q2_OUTPUT_FOLDER, "b")
46     )
47
48     # # Question 3
49     # Q3_OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q3")
50     # if not os.path.exists(Q3_OUTPUT_FOLDER):
51     #     os.makedirs(Q3_OUTPUT_FOLDER)
52     #
53     # q3.learn_binary_factors(
54     #     x=generate_images(),
55     #     k=8,
56     #     em_maximum_iterations=100,
57     #     e_maximum_steps=100,
58     #     e_convergence_criterion=0,
59     # )
```

main.py